

Project_resnet20_4_bit

December 6, 2024

```
[1]: import argparse
import os
import time
import shutil

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.backends.cudnn as cudnn

import torchvision
import torchvision.transforms as transforms

from models import *

global best_prec
use_gpu = torch.cuda.is_available()
print('=> Building model...')

# include resnet model
batch_size = 128
model_name = "project_resnet20"
model = resnet20_quant()

# reduce conv layer input and output size and remove batch norm layer
model.layer1[0].conv1 = QuantConv2d(8, 8, kernel_size=3, stride=1, padding=1,
    ↪ bias=False)
model.layer1[0].bn2 = nn.Sequential()

normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243,
    ↪ 0.262])

train_dataset = torchvision.datasets.CIFAR10(
    root='./data',
```

```

train=True,
download=True,
transform=transforms.Compose([
    transforms.RandomCrop(32, padding=4),
    transforms.RandomHorizontalFlip(),
    transforms.ToTensor(),
    normalize,
]))
trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size,
    ↪shuffle=True, num_workers=2)

test_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        normalize,
    ]))

testloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,
    ↪shuffle=False, num_workers=2)

print_freq = 100 # every 100 batches, accuracy printed. Here, each batch
    ↪includes "batch_size" data points
# CIFAR10 has 50,000 training data, and 10,000 validation data.

def train(trainloader, model, criterion, optimizer, epoch):
    batch_time = AverageMeter()
    data_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    model.train()

    end = time.time()
    for i, (input, target) in enumerate(trainloader):
        # measure data loading time
        data_time.update(time.time() - end)

        input, target = input.cuda(), target.cuda()

        # compute output
        output = model(input)
        loss = criterion(output, target)

```

```

        # measure accuracy and record loss
        prec = accuracy(output, target)[0]
        losses.update(loss.item(), input.size(0))
        top1.update(prec.item(), input.size(0))

        # compute gradient and do SGD step
        optimizer.zero_grad()
        loss.backward()
        optimizer.step()

        # measure elapsed time
        batch_time.update(time.time() - end)
        end = time.time()

    if i % print_freq == 0:
        print('Epoch: [{0}] [{1}/{2}]\t'
              'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
              'Data {data_time.val:.3f} ({data_time.avg:.3f})\t'
              'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
              'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                  epoch, i, len(trainloader), batch_time=batch_time,
                  data_time=data_time, loss=losses, top1=top1))

def validate(val_loader, model, criterion ):
    batch_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to evaluate mode
    model.eval()

    end = time.time()
    with torch.no_grad():
        for i, (input, target) in enumerate(val_loader):

            input, target = input.cuda(), target.cuda()

            # compute output
            output = model(input)
            loss = criterion(output, target)

            # measure accuracy and record loss
            prec = accuracy(output, target)[0]

```

```

        losses.update(loss.item(), input.size(0))
        top1.update(prec.item(), input.size(0))

        # measure elapsed time
        batch_time.update(time.time() - end)
        end = time.time()

        if i % print_freq == 0: # This line shows how frequently print out
            the status. e.g., i%5 => every 5 batch, prints out
                print('Test: [{0}/{1}]\t'
                      'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
                      'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
                      'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                        i, len(val_loader), batch_time=batch_time, loss=losses,
                        top1=top1))

    print(' * Prec {top1.avg:.3f}% '.format(top1=top1))
    return top1.avg

def accuracy(output, target, topk=(1,)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk)
    batch_size = target.size(0)

    _, pred = output.topk(maxk, 1, True, True)
    pred = pred.t()
    correct = pred.eq(target.view(1, -1).expand_as(pred))

    res = []
    for k in topk:
        correct_k = correct[:k].view(-1).float().sum(0)
        res.append(correct_k.mul_(100.0 / batch_size))
    return res

class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

```

```

def update(self, val, n=1):
    self.val = val
    self.sum += val * n
    self.count += n
    self.avg = self.sum / self.count

def save_checkpoint(state, is_best, fdir):
    filepath = os.path.join(fdir, 'checkpoint.pth')
    torch.save(state, filepath)
    if is_best:
        shutil.copyfile(filepath, os.path.join(fdir, 'model_best.pth.tar'))

def adjust_learning_rate(optimizer, epoch):
    adjust_list = [120, 140]
    if epoch in adjust_list:
        for param_group in optimizer.param_groups:
            param_group['lr'] = param_group['lr'] * 0.1

```

=> Building model...

Files already downloaded and verified

Files already downloaded and verified

```

[4]: # inspect the model structure
print(model)

```

```

ResNet_Cifar(
  (conv1): Conv2d(3, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1),
bias=False)
  (bn1): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (layer1): Sequential(
    (0): BasicBlock(
      (conv1): QuantConv2d(
        8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (conv2): QuantConv2d(
        8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (bn1): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (bn2): Sequential()
    )
  )
)

```

```

(1): BasicBlock(
  (conv1): QuantConv2d(
    8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
    (weight_quant): weight_quantize_fn()
  )
  (conv2): QuantConv2d(
    8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
    (weight_quant): weight_quantize_fn()
  )
  (bn1): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (bn2): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
(2): BasicBlock(
  (conv1): QuantConv2d(
    8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
    (weight_quant): weight_quantize_fn()
  )
  (conv2): QuantConv2d(
    8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
    (weight_quant): weight_quantize_fn()
  )
  (bn1): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
  (relu): ReLU(inplace=True)
  (bn2): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
)
)
(layer2): Sequential(
  (0): BasicBlock(
    (conv1): QuantConv2d(
      8, 16, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False
      (weight_quant): weight_quantize_fn()
    )
    (conv2): QuantConv2d(
      16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
      (weight_quant): weight_quantize_fn()
    )
    (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (relu): ReLU(inplace=True)
    (bn2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (downsample): Sequential(
      (0): QuantConv2d(

```

```

        8, 16, kernel_size=(1, 1), stride=(2, 2), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (1): BasicBlock(
        (conv1): QuantConv2d(
            16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (conv2): QuantConv2d(
            16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (bn2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): BasicBlock(
        (conv1): QuantConv2d(
            16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (conv2): QuantConv2d(
            16, 16, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
            (weight_quant): weight_quantize_fn()
        )
        (bn1): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (bn2): BatchNorm2d(16, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    )
    (layer3): Sequential(
        (0): BasicBlock(
            (conv1): QuantConv2d(
                16, 48, kernel_size=(3, 3), stride=(2, 2), padding=(1, 1), bias=False
                (weight_quant): weight_quantize_fn()
            )
            (conv2): QuantConv2d(
                48, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
                (weight_quant): weight_quantize_fn()
            )
        )
    )

```

```

        (bn1): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (relu): ReLU(inplace=True)
        (bn2): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        (downsample): Sequential(
          (0): QuantConv2d(
            16, 48, kernel_size=(1, 1), stride=(2, 2), bias=False
            (weight_quant): weight_quantize_fn()
          )
          (1): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
        )
      )
    (1): BasicBlock(
      (conv1): QuantConv2d(
        48, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (conv2): QuantConv2d(
        48, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (bn1): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (bn2): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
    (2): BasicBlock(
      (conv1): QuantConv2d(
        48, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (conv2): QuantConv2d(
        48, 48, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (bn1): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (relu): ReLU(inplace=True)
      (bn2): BatchNorm2d(48, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    )
  )
  (avgpool): AvgPool2d(kernel_size=8, stride=1, padding=0)
  (fc): Linear(in_features=48, out_features=10, bias=True)
)

```



```
[ ]: # Training session

lr = 3e-2
weight_decay = 1e-4
epochs = 150
best_prec = 0

#model = nn.DataParallel(model).cuda()
model.cuda()
criterion = nn.CrossEntropyLoss().cuda()
optimizer = torch.optim.SGD(model.parameters(), lr=lr, momentum=0.9,
    ↪weight_decay=weight_decay)
#cudnn.benchmark = True

if not os.path.exists('result'):
    os.makedirs('result')
fdir = 'result/'+str(model_name)
if not os.path.exists(fdir):
    os.makedirs(fdir)

for epoch in range(0, epochs):
    adjust_learning_rate(optimizer, epoch)

    train(trainloader, model, criterion, optimizer, epoch)

    # evaluate on test set
    print("Validation starts")
    prec = validate(testloader, model, criterion)

    # remember best precision and save checkpoint
    is_best = prec > best_prec
    best_prec = max(prec, best_prec)
    print('best acc: {:.1f}'.format(best_prec))
    save_checkpoint({
        'epoch': epoch + 1,
        'state_dict': model.state_dict(),
        'best_prec': best_prec,
        'optimizer': optimizer.state_dict(),
    }, is_best, fdir)
```

```
[2]: # Load the checkpoint

PATH = 'result/'+str(model_name)+'/' + 'model_best.pth.tar'
checkpoint = torch.load(PATH)
model.load_state_dict(checkpoint['state_dict'])
device = torch.device("cuda")
```

```

model.cuda()
model.eval()

test_loss = 0
correct = 0

with torch.no_grad():
    for data, target in testloader:
        data, target = data.to(device), target.to(device) # loading to GPU
        output = model(data)
        pred = output.argmax(dim=1, keepdim=True)
        correct += pred.eq(target.view_as(pred)).sum().item()

test_loss /= len(testloader.dataset)

print('\nTest set: Accuracy: {}/{} ({:.0f}%)\n'.format(
    correct, len(testloader.dataset),
    100. * correct / len(testloader.dataset)))

```

Test set: Accuracy: 8853/10000 (89%)

```

[3]: #send an input and grap the value by using prehook like HW3
class SaveOutput:
    def __init__(self):
        self.outputs = []
    def __call__(self, module, module_in):
        self.outputs.append(module_in) # Save the input tensor
    def clear(self):
        self.outputs = []

save_output = SaveOutput()
device = torch.device("cuda" if use_gpu else "cpu")
i = 0

for layer in model.modules():
    i = i+1
    if isinstance(layer, QuantConv2d):
        print(i, "-th layer prehooked")
        layer.register_forward_pre_hook(save_output)

dataiter = iter(trainloader)
images, labels = next(dataiter)
images = images.cuda()
out = model(images)

```

```

7 -th layer prehooked
9 -th layer prehooked
15 -th layer prehooked
17 -th layer prehooked
23 -th layer prehooked
25 -th layer prehooked
32 -th layer prehooked
34 -th layer prehooked
40 -th layer prehooked
44 -th layer prehooked
46 -th layer prehooked
52 -th layer prehooked
54 -th layer prehooked
61 -th layer prehooked
63 -th layer prehooked
69 -th layer prehooked
73 -th layer prehooked
75 -th layer prehooked
81 -th layer prehooked
83 -th layer prehooked

```

```

[6]: # make a integer version to be able to run on the hardware
w_bit = 4

weight_q = model.layer1[0].conv2.weight_q           # quantized value is
↳ stored during the training

w_alpha = model.layer1[0].conv2.weight_quant.wgt_alpha # alpha is defined
↳ in your model already. bring it out here

w_delta = w_alpha / (2**(w_bit-1)-1)                # delta can be
↳ calculated by using alpha and w_bit, for sign number, do minus one for b

weight_int = weight_q / w_delta                      # w_int can be calculated by weight_q and
↳ w_delta

# print(weight_int) # you should see clean integer numbers

```

```

[8]: x_bit = 4
x = save_output.outputs[2][0] # input of the 2nd conv layer
x_alpha = model.layer1[0].conv2.act_alpha
x_delta = x_alpha / (2**x_bit-1)

act_quant_fn = act_quantization(x_bit) # define the quantization function
x_q = act_quant_fn(x, x_alpha)         # create the quantized value for x

x_int = x_q / x_delta
# print(x_int) # you should see clean integer numbers

```

```
[9]: conv_int = torch.nn.Conv2d(in_channels = 8, out_channels=8, kernel_size = 3,
    ↪padding =1, bias = False)
conv_int.weight = torch.nn.parameter.Parameter(weight_int)
relu = nn.ReLU()
bn = nn.BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
    ↪track_running_stats=True).to(device)

output_int = bn(conv_int(x_int))    # output_int can be calculated with
    ↪conv_int and x_int
output_recovered = output_int * x_delta * w_delta # recover with x_delta and
    ↪w_delta
output_recovered = relu(output_recovered)
# print(output_recovered)
```

```
[10]: difference = abs( save_output.outputs[3][0] - output_recovered )
print(difference.mean())    ## It should be small < 1 for the trained model
```

```
tensor(0.2305, device='cuda:0', grad_fn=<MeanBackward0>)
```

```
[ ]: x_int.size()
```

```
[ ]: x_pad = torch.zeros(8, 34, 34).cuda()
x_pad[ :, 1:33, 1:33] = x_int[0].cuda()
X = x_pad[:, 0:2, 0:10]
X.size()
```

```
[ ]: X = torch.reshape(X, (X.size(0), -1))
X.size()
```

```
[ ]: ### storing activation data ###

from pathlib import Path
# Define the folder path
folder_path = Path('./resnet_output/')

# Create the folder if it doesn't exist
folder_path.mkdir(parents=True, exist_ok=True)

bit_precision = 4
file = open('./resnet_output/activation.txt', 'w') #write to file
file.write('#time0row7[msb-lsb],time0row6[msb-lst],...,time0row0[msb-lst]#\n')
file.write('#time1row7[msb-lsb],time1row6[msb-lst],...,time1row0[msb-lst]#\n')
file.write('#.....#\n')

for i in range(X.size(1)): # time step
    for j in range(X.size(0)): # row #
        X_bin = '{0:04b}'.format(round(X[7-j,i].item()))
```

```

        for k in range(bit_precision):
            file.write(X_bin[k])
            #file.write(' ') # for visibility with blank between words, you can use
        file.write('\n')
    file.close() #close file

```

```
[ ]: X[:, 0]
```

```
[ ]: weight_int.size()
```

```
[ ]: W = torch.reshape(weight_int, (weight_int.size(0), weight_int.size(1), -1))
W.size()
```

```

[ ]: ### storing weight data ###

bit_precision = 4

file = open('./resnet_output/weight.txt', 'w')
file.write('#col0row7[msb-lsb],col0row6[msb-lsb],...,col0row0[msb-lsb]#\n')
file.write('#col1row7[msb-lsb],col1row6[msb-lsb],...,col1row0[msb-lsb]#\n')
file.write('#.....#\n')

# be careful about negative numbers
for kij in range(9):
    for i in range(W.size(0)):
        for j in range(W.size(1)): # row #
            if (W[i, 7-j, kij].item() < 0):
                W_bin = '{0:04b}'.format(round(W[i, 7-j, kij].item() +
↪2*bit_precision)) #check again if it works for neg numbers
            else:
                W_bin = '{0:04b}'.format(round(W[i, 7-j, kij].item()))
            for k in range(bit_precision):
                file.write(W_bin[k])
                #file.write(' ') # for visibility with blank between words,
↪you can use
            file.write('\n')
    file.close() #close file

```

```
[ ]: W[0,:,0]
```

```
[ ]: output_int.size()
```

```

[ ]: output = output_int[0]
output = torch.reshape(output, (output.size(0), -1))
O = output[:,0:8]
O.size()

```

```
[ ]: ### Store output value as well ###

bit_precision = 16
file = open('./resnet_output/output.txt', 'w') #write to file
file.write('#time0col7[msb-lsb],time0col6[msb-lsb],...,time0col0[msb-lsb]#\n')
file.write('#time1col7[msb-lsb],time1col6[msb-lsb],...,time1col0[msb-lsb]#\n')
file.write('#.....#\n')

for i in range(O.size(1)): # time step
    for j in range(O.size(0)): # array size
        if (O[7-j,i].item()<0):
            O_bin = '{0:016b}'.format(round(O[7-j,i].item() + 2**bit_precision))
        else:
            O_bin = '{0:016b}'.format(round(O[7-j,i].item()))
        for k in range(bit_precision):
            file.write(O_bin[k])
        #file.write(' ') # for visibility with blank between words, you can use
        file.write('\n')
file.close() #close file
```

[]:

[]:

[]: