

test_custom_VGGNet16

December 6, 2024

```
[2]: import argparse
import os
import time
import shutil

import torch
import torch.nn as nn
import torch.optim as optim
import torch.nn.functional as F
import torch.backends.cudnn as cudnn

import torchvision
import torchvision.transforms as transforms

from models.quant_layer import *
from models.VGG16_custom import *
```

```
[3]: use_gpu = torch.cuda.is_available()
device = torch.device("cuda" if use_gpu else "cpu")
use_gpu, torch.cuda.get_device_name()
```

```
[3]: (True, 'NVIDIA GeForce GTX 1080 Ti')
```

```
[4]: batch_size = 256
model_name = "VGG16_custom1"
model = VGG16_custom()
```

```
[5]: fdir = 'result/'+str(model_name)+'/' + 'model_best.pth.tar'
checkpoint = torch.load(fdir)
model.load_state_dict(checkpoint['state_dict'])
```

```
[5]: <All keys matched successfully>
```

```
[6]: # means and stds for individual RGB channels
# image = (image - mean) / std
normalize = transforms.Normalize(mean=[0.491, 0.482, 0.447], std=[0.247, 0.243, 0.262])
```

```

train_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=True,
    download=True,
    transform=transforms.Compose([
        transforms.RandomCrop(32, padding=4),
        transforms.RandomHorizontalFlip(),
        transforms.ToTensor(),
        normalize,
    ]))

test_dataset = torchvision.datasets.CIFAR10(
    root='./data',
    train=False,
    download=True,
    transform=transforms.Compose([
        transforms.ToTensor(),
        normalize,
    ]))

trainloader = torch.utils.data.DataLoader(train_dataset, batch_size=batch_size,
    ↪shuffle=True, num_workers=2)
testloader = torch.utils.data.DataLoader(test_dataset, batch_size=batch_size,
    ↪shuffle=False, num_workers=2)

```

Files already downloaded and verified
Files already downloaded and verified

```

[7]: print_freq = len(testloader) / 4
    print(print_freq)

```

10.0

```

[8]: def validate(val_loader, model, criterion ):
    batch_time = AverageMeter()
    losses = AverageMeter()
    top1 = AverageMeter()

    # switch to evaluate mode
    model.eval()

    end = time.time()
    with torch.no_grad():
        for i, (input, target) in enumerate(val_loader):

            input, target = input.cuda(), target.cuda()

```

```

    # compute output
    output = model(input)
    loss = criterion(output, target)

    # measure accuracy and record loss
    prec = accuracy(output, target)[0]
    losses.update(loss.item(), input.size(0))
    top1.update(prec.item(), input.size(0))

    # measure elapsed time
    batch_time.update(time.time() - end)
    end = time.time()

    if i % print_freq == 0: # This line shows how frequently print out
        ↪ the status. e.g., i%5 => every 5 batch, prints out
        print('Test: [{0}/{1}]\t'
              'Time {batch_time.val:.3f} ({batch_time.avg:.3f})\t'
              'Loss {loss.val:.4f} ({loss.avg:.4f})\t'
              'Prec {top1.val:.3f}% ({top1.avg:.3f}%)'.format(
                  i, len(val_loader), batch_time=batch_time, loss=losses,
                  top1=top1))

    print(' * Prec {top1.avg:.3f}% '.format(top1=top1))
    return top1.avg

```

```

[9]: def accuracy(output, target, topk=(1,)):
    """Computes the precision@k for the specified values of k"""
    maxk = max(topk)
    batch_size = target.size(0)

    _, pred = output.topk(maxk, 1, True, True) # topk(k, dim=None,
    ↪ largest=True, sorted=True)
                                                # will output (max value, its
    ↪ index)
    pred = pred.t() # transpose
    correct = pred.eq(target.view(1, -1).expand_as(pred)) # "-1": calculate
    ↪ automatically

    res = []
    for k in topk:
        correct_k = correct[:k].view(-1).float().sum(0) # view(-1): make a
        ↪ flattened 1D tensor
        res.append(correct_k.mul_(100.0 / batch_size)) # correct: size of
        ↪ [maxk, batch_size]
    return res

```

```
[10]: class AverageMeter(object):
    """Computes and stores the average and current value"""
    def __init__(self):
        self.reset()

    def reset(self):
        self.val = 0
        self.avg = 0
        self.sum = 0
        self.count = 0

    def update(self, val, n=1):
        self.val = val
        self.sum += val * n    ## n is impact factor
        self.count += n
        self.avg = self.sum / self.count
```

```
[11]: criterion = nn.CrossEntropyLoss().cuda()

model.eval()
model.cuda()

prec = validate(testloader, model, criterion)
```

```
Test: [0/40]    Time 4.897 (4.897)    Loss 0.3328 (0.3328)    Prec 92.578%
(92.578%)
Test: [10/40]   Time 0.044 (0.490)    Loss 0.2382 (0.3521)    Prec 94.141%
(92.045%)
Test: [20/40]   Time 0.044 (0.282)    Loss 0.2607 (0.3790)    Prec 91.406%
(91.574%)
Test: [30/40]   Time 0.029 (0.210)    Loss 0.4853 (0.3782)    Prec 87.891%
(91.557%)
* Prec 91.570%
```

```
[11]: model
```

```
[11]: VGG_quant(
    (features): Sequential(
      (0): QuantConv2d(
        3, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
      )
      (1): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
      (2): ReLU(inplace=True)
      (3): QuantConv2d(
        64, 64, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
```

```

        (weight_quant): weight_quantize_fn()
    )
    (4): BatchNorm2d(64, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (5): ReLU(inplace=True)
    (6): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (7): QuantConv2d(
        64, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (8): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (9): ReLU(inplace=True)
    (10): QuantConv2d(
        128, 128, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (11): BatchNorm2d(128, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (12): ReLU(inplace=True)
    (13): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (14): QuantConv2d(
        128, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (15): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (16): ReLU(inplace=True)
    (17): QuantConv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (18): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (19): ReLU(inplace=True)
    (20): QuantConv2d(
        256, 256, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (21): BatchNorm2d(256, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (22): ReLU(inplace=True)
    (23): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (24): QuantConv2d(

```

```

        256, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (25): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (26): ReLU(inplace=True)
    (27): QuantConv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (28): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (29): ReLU(inplace=True)
    (30): QuantConv2d(
        512, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (31): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (32): ReLU(inplace=True)
    (33): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (34): QuantConv2d(
        512, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (35): BatchNorm2d(8, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (36): ReLU(inplace=True)
    (37): QuantConv2d(
        8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (38): ReLU(inplace=True)
    (39): QuantConv2d(
        8, 512, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
        (weight_quant): weight_quantize_fn()
    )
    (40): BatchNorm2d(512, eps=1e-05, momentum=0.1, affine=True,
track_running_stats=True)
    (41): ReLU(inplace=True)
    (42): MaxPool2d(kernel_size=2, stride=2, padding=0, dilation=1,
ceil_mode=False)
    (43): AvgPool2d(kernel_size=1, stride=1, padding=0)
    )
    (classifier): Linear(in_features=512, out_features=10, bias=True)
)

```

```
[12]: class SaveOutput:
    def __init__(self):
        self.outputs = []
    def __call__(self, module, module_in):
        self.outputs.append(module_in) # Save the input tensor
    def clear(self):
        self.outputs = []

save_output = SaveOutput()
device = torch.device("cuda" if use_gpu else "cpu")
i = 0
count=0

for layer in model.modules():
    i = i+1
    if isinstance(layer, QuantConv2d):
        print(i, "-th layer prehooked")
        layer.register_forward_pre_hook(save_output)
        count = count +1

dataiter = iter(trainloader)
images, labels = next(dataiter)
images = images.cuda()
out = model(images)

print(count)
```

```
3 -th layer prehooked
7 -th layer prehooked
12 -th layer prehooked
16 -th layer prehooked
21 -th layer prehooked
25 -th layer prehooked
29 -th layer prehooked
34 -th layer prehooked
38 -th layer prehooked
42 -th layer prehooked
47 -th layer prehooked
51 -th layer prehooked
54 -th layer prehooked
13
```

```
[13]: layer_input = save_output.outputs[11][0]
layer_output = save_output.outputs[12][0]
layer_input.size(), layer_output.size()
```

```
[13]: (torch.Size([256, 8, 2, 2]), torch.Size([256, 8, 2, 2]))
```

```
[14]: layer_input = layer_input[0]
      layer_output = layer_output[0]
      layer_input.size(), layer_output.size()
```

```
[14]: (torch.Size([8, 2, 2]), torch.Size([8, 2, 2]))
```

```
[15]: # grab data from the 37th layer!!!
```

```
layer = model.features[37]
print(layer)

print(layer._parameters.keys())

print(layer.weight_quant._parameters)
```

```
QuantConv2d(
  8, 8, kernel_size=(3, 3), stride=(1, 1), padding=(1, 1), bias=False
  (weight_quant): weight_quantize_fn()
)
odict_keys(['weight', 'bias', 'act_alpha', 'weight_q'])
OrderedDict([('wgt_alpha', Parameter containing:
tensor(2.3064, device='cuda:0', requires_grad=True))])
```

```
[16]: bw = 4
      weight_q = layer.weight_q
      w_alpha = layer.weight_quant.wgt_alpha
      w_delta = w_alpha / (2**(bw-1)-1)
      w_int = weight_q / w_delta

      print(w_int.shape)
      print(w_int)
```

```
torch.Size([8, 8, 3, 3])
tensor([[[[-0.0000,  0.0000,  0.0000],
          [ 0.0000,  1.0000, -3.0000],
          [-1.0000,  2.0000, -1.0000]],

        [[-1.0000,  4.0000, -2.0000],
          [-3.0000,  7.0000,  2.0000],
          [-0.0000,  5.0000, -4.0000]],

        [[ 1.0000,  7.0000, -5.0000],
          [ 1.0000, -0.0000,  5.0000],
          [-2.0000, -4.0000, -2.0000]],

        [[-1.0000,  2.0000, -1.0000],
          [ 3.0000,  3.0000, -2.0000],
          [-1.0000,  6.0000, -4.0000]],
```



```

[[-1.0000, -5.0000, -5.0000],
 [-2.0000,  1.0000,  0.0000],
 [-2.0000,  0.0000, -1.0000]],

[[-2.0000,  1.0000, -1.0000],
 [ 0.0000, -2.0000, -3.0000],
 [-4.0000, -4.0000,  1.0000]],

[[ 1.0000, -2.0000,  2.0000],
 [ 1.0000,  5.0000,  4.0000],
 [ 0.0000, -1.0000,  2.0000]],

[[-4.0000, -3.0000, -2.0000],
 [-2.0000, -1.0000,  1.0000],
 [-5.0000, -1.0000, -3.0000]],

[[[ 3.0000,  3.0000,  4.0000],
  [-3.0000, -2.0000, -1.0000],
  [-3.0000, -2.0000, -2.0000]],

[[-4.0000,  2.0000,  5.0000],
 [-6.0000,  2.0000, -1.0000],
 [-5.0000,  1.0000,  6.0000]],

[[ 2.0000,  6.0000,  0.0000],
 [-3.0000,  1.0000, -4.0000],
 [-1.0000, -7.0000, -6.0000]],

[[ 1.0000, -4.0000,  1.0000],
 [ 0.0000,  6.0000,  0.0000],
 [-6.0000, -3.0000, -4.0000]],

[[-1.0000,  3.0000,  2.0000],
 [ 6.0000,  7.0000,  0.0000],
 [ 4.0000,  2.0000, -2.0000]],

[[ 3.0000, -6.0000, -1.0000],
 [-2.0000, -4.0000, -7.0000],
 [-3.0000,  2.0000, -4.0000]],

[[-0.0000, -2.0000, -3.0000],
 [ 2.0000,  1.0000,  0.0000],
 [ 1.0000, -0.0000, -2.0000]],

[[-4.0000,  1.0000, -3.0000],
 [-1.0000,  3.0000, -3.0000],

```

```

[-3.0000, 3.0000, 3.0000]]],

[[[-1.0000, -2.0000, -2.0000],
  [-1.0000, 1.0000, -2.0000],
  [-2.0000, -0.0000, 3.0000]],

[[ 1.0000, -3.0000, -2.0000],
 [ 2.0000, -1.0000, 2.0000],
 [ 6.0000, -4.0000, -2.0000]],

[[-3.0000, 0.0000, 4.0000],
 [-2.0000, -2.0000, -1.0000],
 [-1.0000, 7.0000, 1.0000]],

[[-1.0000, -3.0000, -0.0000],
 [-5.0000, 1.0000, 5.0000],
 [ 0.0000, 3.0000, 4.0000]],

[[ 6.0000, 5.0000, -6.0000],
 [ 4.0000, 7.0000, 4.0000],
 [ 0.0000, -1.0000, -1.0000]],

[[ 5.0000, -2.0000, -1.0000],
 [ 1.0000, 3.0000, 2.0000],
 [-1.0000, -1.0000, 2.0000]],

[[ 1.0000, 6.0000, 2.0000],
 [-4.0000, 4.0000, -1.0000],
 [-3.0000, -0.0000, -5.0000]],

[[ 3.0000, -5.0000, -1.0000],
 [ 2.0000, -2.0000, -5.0000],
 [-2.0000, -2.0000, -3.0000]]],

[[[ 0.0000, 4.0000, 0.0000],
  [ 4.0000, 2.0000, -1.0000],
  [-2.0000, 4.0000, 1.0000]],

[[ 2.0000, 3.0000, 3.0000],
 [ 2.0000, -1.0000, -3.0000],
 [-5.0000, -3.0000, -3.0000]],

[[ -0.0000, 6.0000, 3.0000],
 [ 1.0000, 4.0000, -2.0000],
 [-4.0000, 4.0000, 3.0000]],

```

```

[[-2.0000, -4.0000, -1.0000],
 [-2.0000, -4.0000,  5.0000],
 [ 5.0000,  4.0000,  4.0000]],

[[ 3.0000,  4.0000, -2.0000],
 [-5.0000, -1.0000, -4.0000],
 [-7.0000, -0.0000,  0.0000]],

[[ 3.0000,  2.0000,  1.0000],
 [-0.0000,  6.0000,  5.0000],
 [ 2.0000,  1.0000,  0.0000]],

[[-1.0000, -2.0000, -5.0000],
 [-0.0000,  0.0000,  3.0000],
 [-2.0000,  3.0000,  2.0000]],

[[ 4.0000,  0.0000, -4.0000],
 [-2.0000, -5.0000, -2.0000],
 [-0.0000, -4.0000, -3.0000]],

[[[-0.0000, -0.0000, -4.0000],
 [-3.0000, -3.0000, -2.0000],
 [-5.0000, -0.0000, -4.0000]],

[[-0.0000, -2.0000, -3.0000],
 [ 1.0000,  3.0000,  0.0000],
 [-1.0000,  2.0000,  1.0000]],

[[ 0.0000, -6.0000,  5.0000],
 [-2.0000,  5.0000,  5.0000],
 [-2.0000,  0.0000,  1.0000]],

[[ 1.0000,  4.0000, -0.0000],
 [-1.0000,  4.0000, -1.0000],
 [ 3.0000,  5.0000, -5.0000]],

[[-3.0000,  0.0000,  1.0000],
 [ 2.0000, -3.0000,  4.0000],
 [ 1.0000, -2.0000, -3.0000]],

[[ 5.0000,  4.0000,  2.0000],
 [-4.0000, -2.0000,  3.0000],
 [ 3.0000,  3.0000, -5.0000]],

[[ 1.0000,  3.0000, -2.0000],
 [ 1.0000, -1.0000,  5.0000],
 [-6.0000, -5.0000,  0.0000]],

```

```
[[ 3.0000,  5.0000, -2.0000],
 [ 3.0000,  5.0000,  1.0000],
 [ 2.0000, -1.0000, -4.0000]]],
```

```
[[[-3.0000, -3.0000, -2.0000],
 [-3.0000,  5.0000,  4.0000],
 [-4.0000,  1.0000,  1.0000]],
```

```
[[ 2.0000,  2.0000,  1.0000],
 [ 5.0000,  5.0000,  2.0000],
 [-1.0000,  5.0000,  1.0000]],
```

```
[[ -4.0000,  3.0000, -6.0000],
 [-1.0000,  0.0000, -0.0000],
 [-3.0000, -3.0000, -0.0000]],
```

```
[[ -2.0000, -1.0000,  4.0000],
 [-1.0000,  3.0000, -3.0000],
 [-5.0000, -4.0000,  4.0000]],
```

```
[[ 5.0000, -4.0000, -3.0000],
 [ 3.0000, -1.0000, -5.0000],
 [-1.0000,  5.0000, -0.0000]],
```

```
[[ 3.0000, -3.0000,  0.0000],
 [-1.0000,  4.0000,  1.0000],
 [ 1.0000,  5.0000, -1.0000]],
```

```
[[ 1.0000,  0.0000,  1.0000],
 [ 1.0000, -7.0000, -1.0000],
 [-4.0000, -4.0000,  2.0000]],
```

```
[[ -2.0000,  1.0000,  0.0000],
 [-4.0000,  2.0000,  0.0000],
 [-3.0000, -0.0000,  2.0000]]],
```

```
[[[-2.0000, -2.0000,  2.0000],
 [ 2.0000,  1.0000,  2.0000],
 [-1.0000, -1.0000, -3.0000]],
```

```
[[ 1.0000, -0.0000,  4.0000],
 [-2.0000, -3.0000, -5.0000],
 [-1.0000, -1.0000,  2.0000]],
```

```
[[ 0.0000,  6.0000,  1.0000],
```

```

[-1.0000, -0.0000, -1.0000],
[-2.0000, -4.0000, -3.0000]],

[[ 0.0000, -4.0000,  2.0000],
 [-2.0000, -2.0000, -5.0000],
 [-2.0000, -0.0000, -1.0000]],

[[-3.0000,  0.0000,  2.0000],
 [ 3.0000, -7.0000, -3.0000],
 [ 5.0000,  1.0000, -3.0000]],

[[ 0.0000, -1.0000, -4.0000],
 [ 5.0000,  4.0000,  3.0000],
 [ 0.0000,  3.0000,  1.0000]],

[[-4.0000,  2.0000,  2.0000],
 [ 5.0000,  7.0000, -1.0000],
 [ 3.0000,  6.0000, -3.0000]],

[[ 2.0000,  6.0000, -2.0000],
 [ 1.0000,  3.0000,  1.0000],
 [-2.0000,  1.0000, -0.0000]],

[[[-1.0000, -2.0000, -1.0000],
 [ 4.0000,  4.0000,  7.0000],
 [ 1.0000,  2.0000,  1.0000]],

[[ 1.0000, -2.0000,  2.0000],
 [ 4.0000, -0.0000,  7.0000],
 [ 3.0000, -3.0000,  1.0000]],

[[-2.0000,  2.0000, -1.0000],
 [-2.0000, -2.0000, -2.0000],
 [-3.0000, -1.0000, -0.0000]],

[[-0.0000, -1.0000,  2.0000],
 [ 5.0000,  6.0000, -1.0000],
 [ 0.0000, -3.0000,  2.0000]],

[[ 1.0000, -2.0000, -1.0000],
 [ 2.0000, -4.0000, -2.0000],
 [-4.0000,  2.0000, -4.0000]],

[[-4.0000,  2.0000,  0.0000],
 [ 3.0000,  2.0000,  3.0000],
 [-3.0000, -0.0000, -0.0000]],

```

```

[[[-4.0000, -2.0000, -3.0000],
 [ 5.0000,  3.0000, -4.0000],
 [ 3.0000, -2.0000, -2.0000]],

 [[-3.0000,  4.0000, -2.0000],
 [-0.0000,  5.0000,  1.0000],
 [-4.0000, -1.0000,  0.0000]]], device='cuda:0',
grad_fn=<DivBackward0>)

```

```

[17]: x = layer_input
      x_alpha = model.features[37].act_alpha
      x_delta = x_alpha / (2**(bw)-1)

      act_quant_fn = act_quantization(bw)
      x_q = act_quant_fn(x, x_alpha)

      x_int = x_q / x_delta

      print(x_int.shape)
      print(x_int)

```

```

torch.Size([8, 2, 2])
tensor([[[ 0.0000,  0.0000],
          [ 0.0000,  7.0000]],

        [[ 6.0000,  4.0000],
          [ 0.0000,  0.0000]],

        [[ 0.0000,  9.0000],
          [ 0.0000,  0.0000]],

        [[ 0.0000,  0.0000],
          [ 0.0000,  0.0000]],

        [[11.0000,  0.0000],
          [ 1.0000,  0.0000]],

        [[ 6.0000,  2.0000],
          [ 8.0000, 14.0000]],

        [[ 0.0000,  0.0000],
          [ 0.0000,  0.0000]],

        [[ 5.0000,  0.0000],
          [ 0.0000,  0.0000]]], device='cuda:0', grad_fn=<DivBackward0>)

```

```
[18]: conv_int = torch.nn.Conv2d(in_channels=8, out_channels=8, kernel_size=3,
    ↪padding=1, bias=False)
conv_int.weight = torch.nn.parameter.Parameter(w_int)
output_int = F.relu(conv_int(x_int))
output_recovered = output_int * w_delta * x_delta # recover with x_delta and
    ↪w_delta

print(output_recovered.shape)
print(layer_output.shape)
print(output_int)
```

```
torch.Size([8, 2, 2])
torch.Size([8, 2, 2])
tensor([[[ 58.0000,  0.0000],
          [ 0.0000,  9.0000]],

        [[ 0.0000, 16.0000],
          [ 0.0000,  0.0000]],

        [[122.0000, 34.0000],
          [ 71.0000, 162.0000]],

        [[ 0.0000, 42.0000],
          [225.0000, 246.0000]],

        [[ 0.0000, 139.0000],
          [ 83.0000,  0.0000]],

        [[101.0000, 149.0000],
          [ 0.0000, 190.0000]],

        [[ 0.0000, 92.0000],
          [122.0000, 141.0000]],

        [[ 18.0000, 36.0000],
          [100.0000, 74.0000]]], device='cuda:0', grad_fn=<ReluBackward0>)
```

```
[19]: # calculate the difference between outputs, d should be less than 1e-03
diff = abs(layer_output - output_recovered)
print(diff.mean())
```

```
tensor(3.9116e-07, device='cuda:0', grad_fn=<MeanBackward0>)
```

```
[20]: print(x_int.size())
```

```
torch.Size([8, 2, 2])
```

```
[21]: x_pad = torch.zeros(8, 4, 4).cuda()

x_pad[:, 1:3, 1:3] = x_int.cuda()

X = torch.reshape(x_pad, (x_pad.size(0), -1))

print(X.size())
```

torch.Size([8, 16])

```
[22]: from pathlib import Path
# Define the folder path
folder_path = Path('./vgg_output/')

# Create the folder if it doesn't exist
folder_path.mkdir(parents=True, exist_ok=True)
```

```
[23]: ### store weights ###

bit_precision = 4
file = open('./vgg_output/activation.txt', 'w')
file.write('#time0row7[msb-lsb],time0row6[msb-lst],...,time0row0[msb-lst]#\n')
file.write('#time1row7[msb-lsb],time1row6[msb-lst],...,time1row0[msb-lst]#\n')
file.write('#.....#\n')

for i in range(X.size(1)): # time step
    for j in range(X.size(0)): # row #
        X_bin = '{0:04b}'.format(round(X[7-j,i].item()))
        for k in range(bit_precision):
            file.write(X_bin[k])
            #file.write(' ') # use this line for visibility with blank between
↪words
        file.write('\n')
file.close() #close file
```

```
[24]: print(w_int.size())
W = torch.reshape(w_int, (w_int.size(0), w_int.size(1), -1))
W.size()
```

torch.Size([8, 8, 3, 3])

```
[24]: torch.Size([8, 8, 9])
```

```
[25]: bit_precision = 4

for kij in range(9):
    file = open('./vgg_output/w{}.txt'.format(str(kij)), 'w')
```



```

file.write('#col0row7[msb-lsb],col0row6[msb-lsb],...,col0row0[msb-lsb]#\n')
file.write('#col1row7[msb-lsb],col1row6[msb-lsb],...,col1row0[msb-lsb]#\n')
file.write('#.....#\n')
for i in range(W.size(0)):
    for j in range(W.size(1)):
        if (W[i, 7-j, kij].item()<0):
            W_bin = '{0:04b}'.format(int(W[i,7-j,kij].
↪item()+2**bit_precision+0.001))
        else:
            W_bin = '{0:04b}'.format(int(W[i,7-j,kij].item()+0.001))
        for k in range(bit_precision):
            file.write(W_bin[k])
            #file.write(' ') # for visibility with blank between words, you
↪can use
            file.write('\n')
file.close() #close file

```

```

[25]: ### storing weight data ###

bit_precision = 4

file = open('./vgg_output/weight.txt', 'w')
file.write('#col0row7[msb-lsb],col0row6[msb-lsb],...,col0row0[msb-lsb]#\n')
file.write('#col1row7[msb-lsb],col1row6[msb-lsb],...,col1row0[msb-lsb]#\n')
file.write('#.....#\n')

for kij in range(9):
    for i in range(W.size(0)): #col
        for j in range(W.size(1)): # row
            if (W[i, 7-j, kij].item()<0):
                W_bin = '{0:04b}'.format(round(W[i,7-j, kij].item() +
↪2**bit_precision)) #check again if it works for neg numbers
            else:
                W_bin = '{0:04b}'.format(round(W[i,7-j, kij].item()))
            for k in range(bit_precision):
                file.write(W_bin[k])
                #file.write(' ') # for visibility with blank between words,
↪you can use
                file.write('\n')
file.close() #close file

```

```

[26]: print(output_int.size())
0 = torch.reshape(output_int, (output_int.size(0), -1))
print(0.size())

```

```

torch.Size([8, 2, 2])
torch.Size([8, 4])

```

```
[27]: ### Store output data ###

bit_precision = 16
file = open('./vgg_output/output.txt', 'w') #write to file
file.write('#time0col7[msb-lsb],time0col6[msb-lsb],...,time0col0[msb-lsb]#\n')
file.write('#time1col7[msb-lsb],time1col6[msb-lsb],...,time1col0[msb-lsb]#\n')
file.write('#.....#\n')

for i in range(0.size(1)):
    for j in range(0.size(0)):
        if (0[7-j,i].item()<0):
            O_bin = '{0:016b}'.format(round(0[7-j,i].item() + 2**bit_precision))
        else:
            O_bin = '{0:016b}'.format(round(0[7-j,i].item()))
        for k in range(bit_precision):
            file.write(O_bin[k])
        #file.write(' ') # for visibility with blank between words, you can use
        file.write('\n')
file.close() #close file
```

```
[28]: print(X.size())
```

```
torch.Size([8, 16])
```

```
[29]: psum = torch.zeros(8, 16, 9).cuda() #initialize an empty psum first with array_
      ↪ size, p_nij and kij
print(psum.size())

# calculate psum value
for kij in range(9):
    for p_nij in range(16):
        m = nn.Linear(8, 8, bias=False) # array size matched
        m.weight = torch.nn.Parameter(W[:, :, kij])
        psum[:, p_nij, kij] = m(X[:, p_nij]).cuda()
```

```
torch.Size([8, 16, 9])
```

```
[30]: ### Store psum data, not needed for output stationary model ###

bit_precision = 16

file = open('./vgg_output/psum.txt', 'w') #write to file
file.write('#time0col7[msb-lsb],time0col6[msb-lsb],...,time0col0[msb-lsb]#\n')
file.write('#time1col7[msb-lsb],time1col6[msb-lsb],...,time1col0[msb-lsb]#\n')
file.write('#.....#\n')

for kij in range(9):
```

```

for i in range(psum.size(1)): # time step
    for j in range(psum.size(0)): # array size
        if (psum[7-j,i, kij].item()<0):
            psum_bin = '{0:016b}'.format(round(psum[7-j,i, kij].item() +
↪2**bit_precision))
        else:
            psum_bin = '{0:016b}'.format(round(psum[7-j,i, kij].item()))
        for k in range(bit_precision):
            file.write(psum_bin[k])
            #file.write(' ') # for visibility with blank between words, you
↪can use
            file.write('\n')
file.close() #close file

```

[]:

[]: